

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv("superstore.csv", encoding='iso-8859-1')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 51290 entries, 0 to 51289
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Row ID                51290 non-null  int64
1   Order ID              51290 non-null  object
2   Order Date            51290 non-null  object
3   Ship Date             51290 non-null  object
4   Ship Mode             51290 non-null  object
5   Customer ID           51290 non-null  object
6   Customer Name         51290 non-null  object
7   Segment              51290 non-null  object
8   City                  51290 non-null  object
9   State                 51290 non-null  object
10  Country               51290 non-null  object
11  Postal Code           9994 non-null   float64
12  Market               51290 non-null  object
13  Region               51290 non-null  object
14  Product ID           51290 non-null  object
15  Category             51290 non-null  object
16  Sub-Category         51290 non-null  object
17  Product Name         51290 non-null  object
18  Sales                51290 non-null  float64
19  Quantity             51290 non-null  int64
20  Discount             51290 non-null  float64
21  Profit               51290 non-null  float64
22  Shipping Cost        51290 non-null  float64
23  Order Priority        51290 non-null  object
dtypes: float64(5), int64(2), object(17)
memory usage: 9.4+ MB
```

```
In [3]: df = df.drop('Postal Code', axis=1)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 51290 entries, 0 to 51289
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Row ID                51290 non-null  int64
1   Order ID              51290 non-null  object
2   Order Date            51290 non-null  object
3   Ship Date             51290 non-null  object
4   Ship Mode             51290 non-null  object
5   Customer ID           51290 non-null  object
6   Customer Name         51290 non-null  object
7   Segment              51290 non-null  object
8   City                 51290 non-null  object
9   State                51290 non-null  object
10  Country              51290 non-null  object
11  Market              51290 non-null  object
12  Region              51290 non-null  object
13  Product ID          51290 non-null  object
14  Category            51290 non-null  object
15  Sub-Category        51290 non-null  object
16  Product Name        51290 non-null  object
17  Sales               51290 non-null  float64
18  Quantity            51290 non-null  int64
19  Discount            51290 non-null  float64
20  Profit              51290 non-null  float64
21  Shipping Cost       51290 non-null  float64
22  Order Priority       51290 non-null  object
dtypes: float64(4), int64(2), object(17)
memory usage: 9.0+ MB
```

```
In [5]: df['Order Date'] = pd.to_datetime(df['Order Date'])
df['Ship Date'] = pd.to_datetime(df['Ship Date'])
```

```
In [6]: df.dtypes
```


```
Out[6]: Row ID          int64
Order ID         object
Order Date       datetime64[ns]
Ship Date        datetime64[ns]
Ship Mode        object
Customer ID      object
Customer Name    object
Segment          object
City             object
State            object
Country          object
Market           object
Region           object
Product ID       object
Category         object
Sub-Category     object
Product Name     object
Sales            float64
Quantity         int64
Discount         float64
Profit           float64
Shipping Cost    float64
Order Priority    object
dtype: object
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: Row ID          0
Order ID         0
Order Date       0
Ship Date        0
Ship Mode        0
Customer ID      0
Customer Name    0
Segment          0
City             0
State            0
Country          0
Market           0
Region           0
Product ID       0
Category         0
Sub-Category     0
Product Name     0
Sales            0
Quantity         0
Discount         0
Profit           0
Shipping Cost    0
Order Priority    0
dtype: int64
```

```
In [8]: plt.figure(figsize = (25, 25))
sns.heatmap(df.isnull())
```

```
Out[8]: <AxesSubplot:>
```



```
In [9]: df.select_dtypes(include=['int64', 'float64']).columns
```

```
Out[9]: Index(['Row ID', 'Sales', 'Quantity', 'Discount', 'Profit', 'Shipping Cost'], dtype='object')
```

```
In [10]: df.select_dtypes(include=['object']).columns
```

```
Out[10]: Index(['Order ID', 'Ship Mode', 'Customer ID', 'Customer Name', 'Segment',
               'City', 'State', 'Country', 'Market', 'Region', 'Product ID',
               'Category', 'Sub-Category', 'Product Name', 'Order Priority'],
               dtype='object')
```

```
In [11]: df.select_dtypes(include=['datetime64[ns]']).columns
```

```
Out[11]: Index(['Order Date', 'Ship Date'], dtype='object')
```

# Basic Exploratory data analysis

What are the total sales and profits for the store for each year in the datasets?

```
In [12]: # What are the total sales and profits for the store for each year in the datasets?

# create new columns (year)
df['year'] = df['Order Date'].dt.year
#df.head()
#Calculate the total sales for each year?
sales_by_year = df.groupby('year')['Sales'].sum()
#sales_by_year
#calculate the total shipping cost of the year?
shipping_cost_year = df.groupby('year')['Shipping Cost'].sum()
#shipping_cost_year

# calculate the total profits for each year?
profits_by_year = sales_by_year - shipping_cost_year
profits_by_year
```

```
Out[12]: year
2011      2.015180e+06
2012      2.393947e+06
2013      3.041196e+06
2014      3.839358e+06
dtype: float64
```

What are the top-selling products in terms of revenue and quantity sold?

```
In [13]: # What are the top-selling products in terms of revenue and quantity sold?
product_sales = df.groupby('Product Name').agg({'Sales': 'sum',
                                                'Quantity': 'sum'})

#product_sales.sort_values(by='Sales', ascending=False).head(10)
product_sales.sort_values(by='Quantity', ascending=False).head(10)
```

Out[13]:

	Sales	Quantity
Product Name		
Staples	7008.2000	876
Cardinal Index Tab, Clear	1922.8302	337
Eldon File Cart, Single Width	34387.7287	321
Rogers File Cart, Single Width	29466.3053	262
Sanford Pencil Sharpener, Water Color	5581.9741	259
Stockwell Paper Clips, Assorted Sizes	2395.4378	253
Avery Index Tab, Clear	1312.9488	252
Ibico Index Tab, Clear	1807.1811	251
Smead File Cart, Single Width	25397.1708	250
Stanley Pencil Sharpener, Water Color	5537.2290	242

## What are the top product categories in terms of revenue and quantity sold?

```
In [14]: #What are the top product categories in terms of revenue and quantity sold?
category_sales = df.groupby('Category').agg({'Sales': 'sum',
                                             'Quantity': 'sum'})
#category_sales.sort_values(by='Sales', ascending=False).head(10)
category_sales.sort_values(by='Quantity', ascending=False).head(10)
```

Out[14]:

	Sales	Quantity
Category		
Office Supplies	3.787070e+06	108182
Technology	4.744557e+06	35176
Furniture	4.110874e+06	34954

## what is the average order value and how does it vary by product category?

```
In [15]: # what is the average order value and how does it vary by product category?
order_revenue = df.groupby('Order ID')['Sales'].sum()
avg_order_value = order_revenue.mean()

# calculated the avg order value by product category
category_avg_order_value = df.groupby('Category')['Sales'].mean()

print("Average Order Value: ${:.2f}".format(avg_order_value))
print("\nAverage Order Value by Product Category: ")
print(category_avg_order_value)
```

Average Order Value: \$504.99

Average Order Value by Product Category:

Category	
Furniture	416.248905
Office Supplies	121.097120
Technology	467.858939

Name: Sales, dtype: float64

## Who are the top customers in terms of total spending and number of orders?

```
In [16]: #Who are the top customers in terms of total spending and number of orders?

customer_spending = df.groupby("Customer Name")["Sales"].sum()
top_customers_by_spending = customer_spending.sort_values(ascending=False).head(3)
customer_orders = df.groupby("Customer Name")["Order ID"].nunique()
top_customers_by_orders = customer_orders.sort_values(ascending=False).head(3)

print("Top 3 Customers by total spending:")
print(top_customers_by_spending)
print("\nTop 3 Customers by Number of Orders:")
print(top_customers_by_orders)
```

Top 3 Customers by total spending:

Customer Name	
Tom Ashbrook	40488.07080
Tamara Chand	37457.33300
Greg Tran	35550.95428

Name: Sales, dtype: float64

Top 3 Customers by Number of Orders:

Customer Name	
Frank Olsen	47
Anna Andreadi	47
Michael Paige	47

Name: Order ID, dtype: int64

## How does sales performance vary across different regions or states?

```
In [17]: #How does sales performance vary across different regions or states?

region_sales = df.groupby("Region")["Sales"].sum()
state_sales = df.groupby('State')['Sales'].sum()

top_region_sales = region_sales.sort_values(ascending=False)
top_state_sales = state_sales.sort_values(ascending=False).head(10)

print("Total sales by region:")
print(top_region_sales)
print("\nTotal sales by states:")
print(top_state_sales)
```

Total sales by region:

Region	
Central	2.822303e+06
South	1.600907e+06
North	1.248166e+06
Oceania	1.100185e+06
Southeast Asia	8.844232e+05
North Asia	8.483098e+05
EMEA	8.061613e+05
Africa	7.837732e+05
Central Asia	7.528266e+05
West	7.254578e+05
East	6.787812e+05
Caribbean	3.242809e+05
Canada	6.692817e+04

Name: Sales, dtype: float64

Total sales by states:

State	
England	485170.9710
California	457687.6315
Ile-de-France	317822.5440
New York	310876.2710
New South Wales	270487.1040
Queensland	238312.7340
North Rhine-Westphalia	216451.8510
Texas	170188.0458
San Salvador	153639.3970
National Capital	152175.3555


Name: Sales, dtype: float64

## How do sales trends vary across different time periods (e.g., months, quarters, years)?

```
In [18]: #How do sales trends vary across different time periods (e.g., months, quarters, years)
monthly_sales = df.groupby(pd.Grouper(key='Order Date', freq='M'))['Sales'].sum()
quarterly_sales = df.groupby(pd.Grouper(key='Order Date', freq='Q'))['Sales'].sum()
yearly_sales = df.groupby(pd.Grouper(key='Order Date', freq='Y'))['Sales'].sum()

#plot the sales trends over time
plt.plot(monthly_sales, label='Monthly Sales')
plt.plot(quarterly_sales, label = 'Quarterly sales')
plt.plot(yearly_sales, label = "Yearly Sales")
plt.xticks(rotation = 50)
plt.xlabel('Time')
plt.ylabel("Sales")
plt.legend()
plt.show()
```





Can we use the data to forecast future sales and identify potential growth opportunities for the business?

```
In [20]: from statsmodels.tsa.arima.model import ARIMA
```


```
In [21]: # Create a new DataFrame for sales data by month
monthly_sales = df.groupby(pd.Grouper(key='Order Date', freq='M')).sum()

# Fit an ARIMA model to the monthly sales data
model = ARIMA(monthly_sales['Sales'], order=(1, 1, 1))
model_fit = model.fit()

# Generate forecasted sales data for the next 12 months
forecast = model_fit.forecast(steps=12)

# Plot the historical sales data and forecasted sales data
plt.plot(monthly_sales.index, monthly_sales['Sales'])
plt.plot(forecast.index, forecast, color='red')
plt.title('Historical Sales and Forecasted Sales')
plt.xlabel('Month')
plt.ylabel('Sales')

plt.show()
```



## Data visualization

```
In [22]: # Top customers by sales

# group the data by customers name and location and calculated total sales

customer_sales = df.groupby(['Customer Name', 'Segment', 'City', 'State']).agg({
    'Sales': 'sum'
}).reset_index()

#print(customer_sales)
# sort results by sales

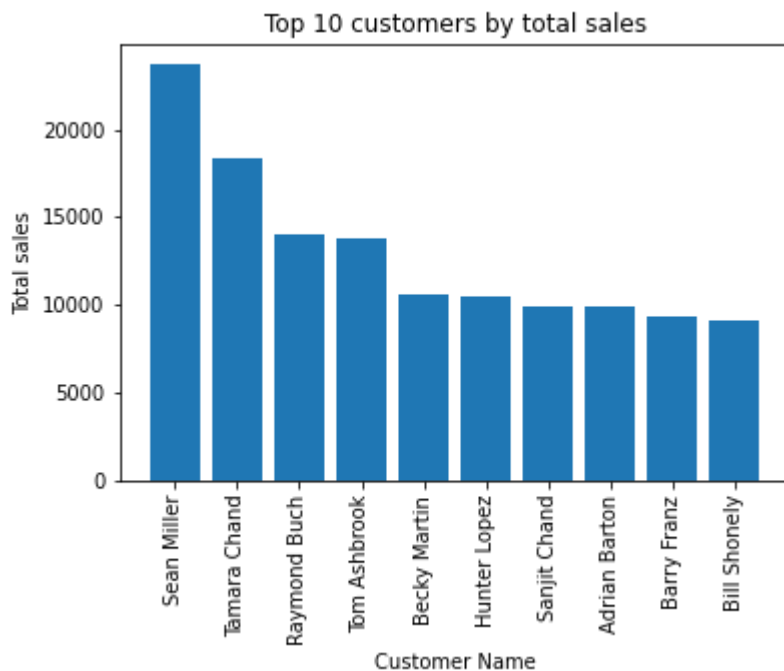
customer_sales = customer_sales.sort_values(by = 'Sales', ascending = False)
#print(customer_sales)

# Creating bar chart showing top 10 customers by total sales

top_10 = customer_sales.head(10)
plt.bar(top_10['Customer Name'], top_10['Sales'])
plt.xticks(rotation=90) # for x axis 90 degree
plt.xlabel("Customer Name")
plt.ylabel("Total sales")
plt.title("Top 10 customers by total sales")
plt.show()

# Display the order history for the top 10 customers

top_10_history = df[df['Customer Name'].isin(top_10['Customer Name'])]
print(top_10_history[['Customer Name', 'Order Date', 'Region', 'Category', 'Sales']])
```




	Customer Name	Order Date	Region	Category	Sales
336	Becky Martin	2011-01-12	South	Office Supplies	57.12
339	Becky Martin	2011-01-12	South	Office Supplies	38.68
393	Hunter Lopez	2011-02-04	Central Asia	Technology	1278.00
397	Hunter Lopez	2011-02-04	Central Asia	Technology	3195.00
404	Hunter Lopez	2011-02-04	Central Asia	Office Supplies	149.58
...	...	...	...	...	...
51237	Bill Shonely	2014-12-31	Central Asia	Office Supplies	258.12
51245	Bill Shonely	2014-12-31	Central Asia	Technology	276.60
51269	Bill Shonely	2014-12-31	Central Asia	Furniture	364.59
51271	Bill Shonely	2014-12-31	Central Asia	Office Supplies	72.00
51276	Bill Shonely	2014-12-31	Central Asia	Office Supplies	39.42

[703 rows x 5 columns]


```
In [30]: # Group the data by quarter and region or category
sales_by_quarter = df.groupby([pd.Grouper(key='Order Date', freq='Q'), 'Region'])['Sales']
# Or, we can group the data by quarter and category
# sales_by_quarter = data.groupby([pd.Grouper(key='Order Date', freq='Q'), 'Category'])

# Plot the line chart
fig, ax = plt.subplots(figsize=(10, 6))
for region in sales_by_quarter['Region'].unique():
    ax.plot(sales_by_quarter[sales_by_quarter['Region']==region]['Order Date'],
            sales_by_quarter[sales_by_quarter['Region']==region]['Sales'],
            label=region)
plt.xticks(rotation=90)
ax.set_xlabel('Quarter')
ax.set_ylabel('Total Sales')
ax.set_title('Sales Trends by Quarter')
ax.legend()
plt.show()
```



```
In [23]: # Get monthly sales data
monthly_sales = df.groupby(pd.Grouper(key='Order Date', freq='M')).agg({'Sales': 'sum'})

# Plot monthly sales trend
plt.plot(monthly_sales.index, monthly_sales['Sales'])
plt.xticks(rotation=50)
plt.title("Monthly Sales Trend")
plt.xlabel("Month")
plt.ylabel("Sales")
plt.show()
```




```
In [24]: #quarterly sales trend
quarterly_sales = df.groupby(pd.Grouper(key='Order Date', freq='Q')).agg({'Sales': 'sum'})
```

```
# Plot quarterly sales trend
plt.plot(quarterly_sales.index, quarterly_sales['Sales'], color = 'green')
plt.xticks(rotation=50)
plt.title("Quarterly Sales Trend")
plt.xlabel("Quarter")
plt.ylabel("Sales")
plt.show()
```




```
In [25]: # yearly sales trend
yearly_sales = df.groupby(pd.Grouper(key='Order Date', freq='Y')).agg({'Sales': 'sum'})
# Plot yearly sales trend
plt.plot(yearly_sales.index, yearly_sales['Sales'], color = 'yellow')
plt.xticks(rotation=50)
plt.title("Yearly Sales Trend")
plt.xlabel("Year")
plt.ylabel("Sales")
plt.show()
```



```
In [26]: # Get top 10 products by revenue
top_products = df.groupby(['Product Name'])['Sales'].sum().reset_index().sort_values('Sales', ascending=False)
```

```
# Plot top 10 products by revenue
plt.bar(top_products['Product Name'], top_products['Sales'])
plt.xticks(rotation=90)
plt.title("Top 10 Products by Revenue")
plt.xlabel("Product Name")
plt.ylabel("Revenue")
plt.show()
```



```
In [27]: # Calculate profit margin for each sale
df["Profit Margin"] = (df["Profit"] / df["Sales"]).round(2)


# Group the data by category and region and sum the profit margins
profit_margin_df = df.groupby(["Category", "Region"])["Profit Margin"].sum().reset_index()

# Pivot the data to create a table with columns for each region
profit_margin_pivot = profit_margin_df.pivot(index="Category", columns="Region", values="Profit Margin")

# Create a stacked bar chart
plt.figure(figsize = (25, 25))
profit_margin_pivot.plot(kind="bar", stacked=True)
plt.title("Profit Margins by Category and Region")
plt.xlabel("Category")
```

```
plt.ylabel("Total Profit Margin")
plt.show()
```


<Figure size 1800x1800 with 0 Axes>



```
In [28]: # Step 4: Add a new column 'Day of Week' to the dataframe
df['Day of Week'] = df['Order Date'].dt.day_name()
# Step 5: Group the dataframe by 'Day of Week', and aggregate the 'Sales' column
sales_by_day = df.groupby('Day of Week')['Sales'].sum()

# Step 6: Filter the data by category or region, if desired
# Example:
# sales_by_day = df[df['Category'] == 'Office Supplies'].groupby('Day of Week')['Sales']

# Step 7: Create a line chart using matplotlib
plt.plot(sales_by_day.index, sales_by_day.values)
plt.xticks(rotation = 50)
plt.title('Total Sales by Day of Week')
plt.xlabel('Day of Week')
plt.ylabel('Sales')
plt.show()
```




```
In [55]: #Aggregate the data by product and customer segment
grouped = df.groupby(['Product Name', 'Segment'], as_index=False)['Sales'].sum()

# Create a scatter plot
fig, ax = plt.subplots(figsize=(10, 6))
colors = {'Consumer': 'red', 'Corporate': 'blue', 'Home Office': 'green'}
for segment, color in colors.items():
    x = grouped[grouped['Segment'] == segment]['Product Name']
    y = grouped[grouped['Segment'] == segment]['Sales']
    ax.scatter(x, y, color=color, label=segment)
ax.legend()
ax.set_xlabel('Product Name')
ax.set_ylabel('Total Sales')
ax.set_title('Sales by Product and Customer Segment')
plt.xticks(rotation=90)
plt.show()
```

```
C:\Users\ASUS\AppData\Local\Programs\Python\Python310\lib\site-packages\IPython\core
\pylabtools.py:151: UserWarning: Glyph 148 (\x94) missing from current font.
  fig.canvas.print_figure(bytes_io, **kw)
C:\Users\ASUS\AppData\Local\Programs\Python\Python310\lib\site-packages\IPython\core
\pylabtools.py:151: UserWarning: Glyph 147 (\x93) missing from current font.
  fig.canvas.print_figure(bytes_io, **kw)
```



### Sales by Product and Customer Segment

[illegible]


```
In [29]: # convert the order date to a datetime object
df['Order Date'] = pd.to_datetime(df['Order Date'])

# group the data by year and customer name
grouped = df.groupby([df['Order Date'].dt.year, 'Customer Name'])

# calculate the number of orders and number of repeat customers
orders = grouped.size()
repeat_customers = (orders > 1).groupby(level=0).sum()

# calculate the retention rate
retention_rate = repeat_customers / orders.groupby(level=0).size()

# plot the retention rate over time
retention_rate.plot()
plt.xlabel('Year')
plt.ylabel('Retention Rate')
plt.title('Customer Retention Rates by Year')
plt.show()
```



## Statistical measurement to better understand data

```
In [30]: # mean
mean_sales = df['Sales'].mean()
#median
median_sales = df['Sales'].median()
#mode
mode_sales = df['Sales'].mode()[0]

print("MeanSales: {:.2f}".format(mean_sales))
print("MedianSales: {:.2f}".format(median_sales))
print("ModeSales: {:.2f}".format(mode_sales))
```

MeanSales: \$246.49  
MedianSales: \$85.05  
ModeSales: \$12.96

```
In [31]: # Range of sales
sales_range = df['Sales'].max()-df['Sales'].min()
print("Sales range: ${:.2f}".format(sales_range))
```

Sales range: \$22638.04

```
In [32]: total_sales = df['Sales'].sum()
total_profit = df['Profit'].sum()
print("Total Sales: ${:.2f}".format(total_sales))
print("Total Profit: ${:.2f}".format(total_profit))
```

Total Sales: \$12642501.91  
Total Profit: \$1467457.29

```
In [33]: order_counts = df['Region'].value_counts()
print("Total order: ", order_counts)
```


Total order: Central 11117  
South 6645  
EMEA 5029  
North 4785  
Africa 4587  
Oceania 3487  
West 3203  
Southeast Asia 3129  
East 2848  
North Asia 2338  
Central Asia 2048  
Caribbean 1690  
Canada 384  
Name: Region, dtype: int64

```
In [34]: sales_profit_corr = df['Sales'].corr(df['Profit'])
print(sales_profit_corr)
```

0.4849181126194441

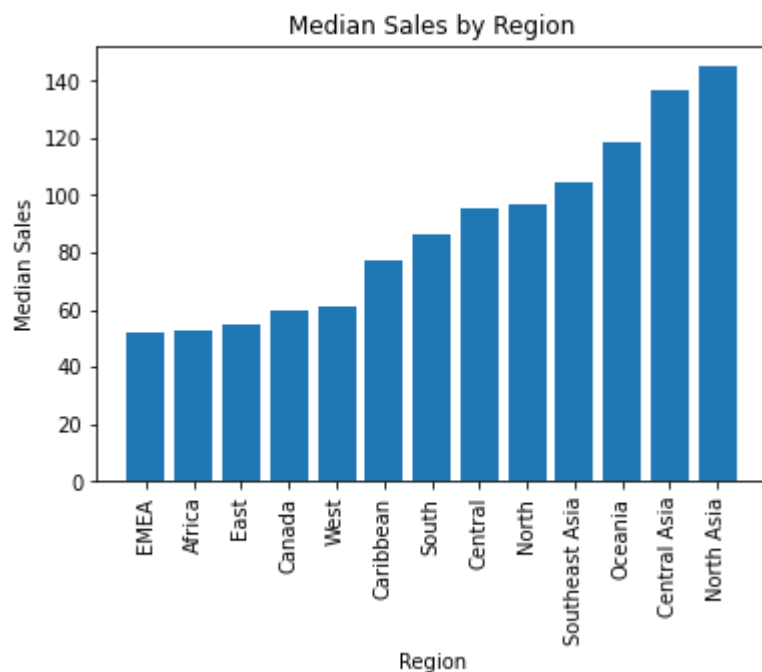
```
In [35]: mean_sales_by_category = df.groupby('Category')['Sales'].mean().sort_values()

plt.bar(mean_sales_by_category.index, mean_sales_by_category.values)
plt.xticks(rotation=90)
plt.xlabel('Category')
plt.ylabel('Mean Sales')
plt.title('Mean Sales by Category')
plt.show()
```



```
In [36]: median_sales_by_region = df.groupby('Region')['Sales'].median().sort_values()


plt.bar(median_sales_by_region.index, median_sales_by_region.values)
plt.xticks(rotation=90)
plt.xlabel('Region')
plt.ylabel('Median Sales')
plt.title('Median Sales by Region')
plt.show()
```



```
In [37]: std_profits_by_category = df.groupby('Category')['Profit'].std().sort_values()

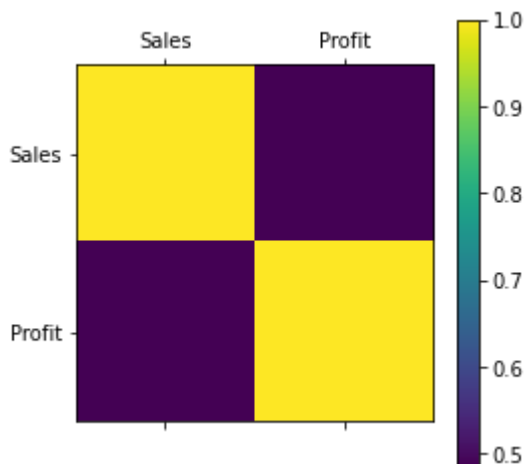
plt.bar(std_profits_by_category.index, std_profits_by_category.values)
plt.xticks(rotation=90)
plt.xlabel('Category')
plt.ylabel('Standard Deviation of Profits')
```

```
plt.title('Standard Deviation of Profits by Category')
plt.show()
```



```
In [38]: sales_and_profits = df[['Sales', 'Profit']]
corr = sales_and_profits.corr()


plt.matshow(corr)
plt.xticks(range(len(sales_and_profits.columns)), sales_and_profits.columns)
plt.yticks(range(len(sales_and_profits.columns)), sales_and_profits.columns)
plt.colorbar()
plt.show()
```



```
In [39]: # Calculate the probability distribution of Sales
sales_prob = df['Sales'].value_counts(normalize=True)

# Plot the probability distribution using a histogram
sns.histplot(df, x="Sales", stat="probability", bins=30)
```


```
Out[39]: <AxesSubplot:xlabel='Sales', ylabel='Probability'>
```



```
In [40]: # Calculate the correlation matrix
corr_matrix = df.corr()

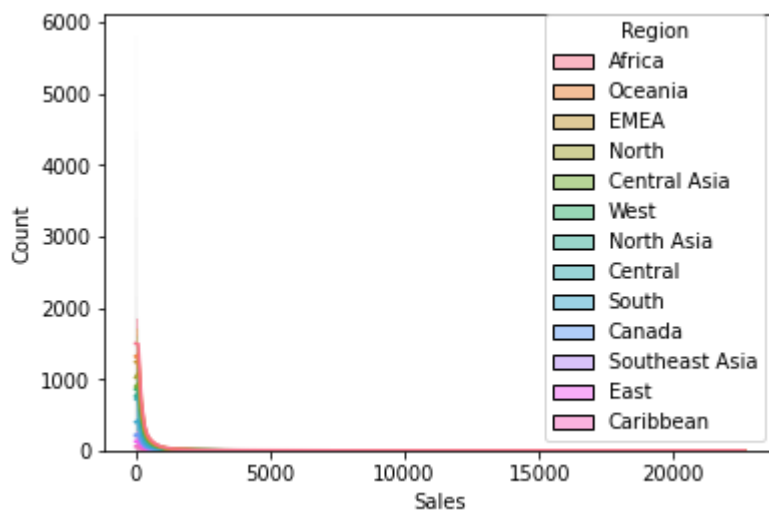
# Visualize the correlation matrix as a heatmap
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
```

Out[40]: <AxesSubplot:>




```
In [41]: # Use Seaborn to create a histogram of sales by region
sns.histplot(data=df, x='Sales', hue='Region', kde=True, multiple='stack')
```


Out[41]: <AxesSubplot:xlabel='Sales', ylabel='Count'>




```
In [42]: sns.scatterplot(x='Sales', y='Profit', data=df)
plt.title('Sales vs. Profit')
plt.show()
```



```
In [78]: sns.histplot(x='Sales', data=df)
plt.title('Distribution of Sales')
plt.show()
```




```
In [43]: sns.barplot(x='Region', y='Sales', data=df)
plt.xticks(rotation = 90)
plt.title('Sales by Region')
plt.show()
```

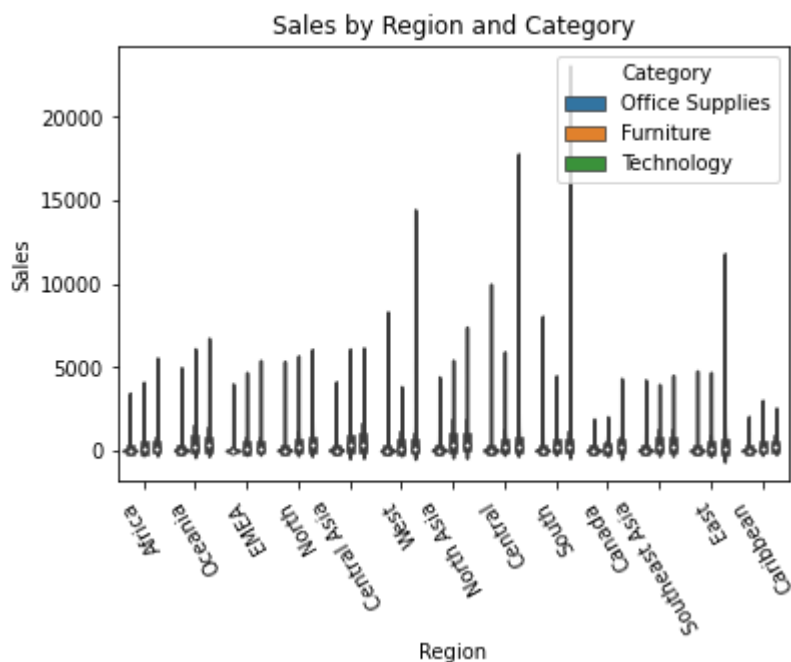


```
In [82]: sns.boxplot(x='Category', y='Sales', data=df)
plt.title('Sales by Category')
plt.show()
```







```
In [85]: sns.violinplot(x='Region', y='Sales', hue='Category', data=df)
plt.xticks(rotation = 120)
plt.title('Sales by Region and Category')
plt.show()
```



```
In [86]: sns.countplot(x='Ship Mode', hue='Region', data=df)
plt.title('Ship Mode by Region')
plt.show()
```



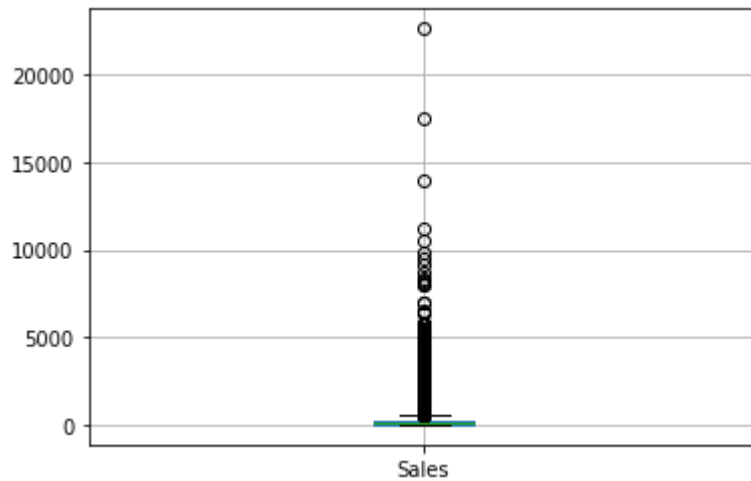
```
In [87]: df['Month'] = pd.to_datetime(df['Order Date']).dt.month
sales_by_month = df.groupby('Month')['Sales'].sum().reset_index()
sns.lineplot(x='Month', y='Sales', data=sales_by_month)
plt.title('Sales by Month')
plt.show()
```



```
In [44]: # Identify outliers using box plot
df.boxplot(column=['Sales'])

# Calculate upper and lower bounds
q1 = df['Sales'].quantile(0.25)
q3 = df['Sales'].quantile(0.75)
iqr = q3 - q1
upper_bound = q3 + 1.5 * iqr
lower_bound = q1 - 1.5 * iqr

# Remove outliers
df = df[(df['Sales'] > lower_bound) & (df['Sales'] < upper_bound)]
```



```
In [45]: from sklearn.preprocessing import LabelEncoder
```

```
In [46]: # Label encode the Category column
le = LabelEncoder()
df['Category'] = le.fit_transform(df['Category'])

# One-hot encode the Region column
df = pd.get_dummies(df, columns=['Region'])
```

C:\Users\ASUS\AppData\Local\Temp\ipykernel\_9848\1845053093.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
df['Category'] = le.fit\_transform(df['Category'])

```
In [48]: from sklearn.preprocessing import StandardScaler
```

```
In [49]: from sklearn.preprocessing import MinMaxScaler
```

```
In [50]: # Select the numerical columns to be scaled
numerical_columns = ['Sales', 'Quantity', 'Discount']

# Scale the numerical columns using MinMaxScaler
scaler = MinMaxScaler()
df[numerical_columns] = scaler.fit_transform(df[numerical_columns])

# View the scaled dataframe
print(df.head())
```

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	\
0	42433	AG-2011-2040	2011-01-01	2011-06-01	Standard Class	TB-11280	
1	22253	IN-2011-47883	2011-01-01	2011-08-01	Standard Class	JH-15985	
2	48883	HU-2011-1220	2011-01-01	2011-05-01	Second Class	AT-735	
3	11731	IT-2011-3647632	2011-01-01	2011-05-01	Second Class	EM-14140	
4	22255	IN-2011-47883	2011-01-01	2011-08-01	Standard Class	JH-15985	

	Customer Name	Segment	City	State	...	\
0	Toby Braunhardt	Consumer	Constantine	Constantine	...	
1	Joseph Holt	Consumer	Wagga Wagga	New South Wales	...	
2	Annie Thurman	Consumer	Budapest	Budapest	...	
3	Eugene Moren	Home Office	Stockholm	Stockholm	...	
4	Joseph Holt	Consumer	Wagga Wagga	New South Wales	...	

	Region_Central	Region_Central	Asia	Region_EMEA	Region_East	Region_North	\
0	0		0	0	0	0	
1	0		0	0	0	0	
2	0		0	1	0	0	
3	0		0	0	0	1	
4	0		0	0	0	0	

	Region_North	Asia	Region_Oceania	Region_South	Region_Southeast	Asia	\
0	0		0	0		0	
1	0		1	0		0	
2	0		0	0		0	
3	0		0	0		0	
4	0		1	0		0	

	Region_West
0	0
1	0
2	0
3	0
4	0

[5 rows x 38 columns]

```
In [51]: # Create a new feature "profit" by subtracting the cost from the sale price
df['profit'] = df['Sales'] - df['Profit']

# Display the updated dataset
print(df.head())
```

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	\
0	42433	AG-2011-2040	2011-01-01	2011-06-01	Standard Class	TB-11280	
1	22253	IN-2011-47883	2011-01-01	2011-08-01	Standard Class	JH-15985	
2	48883	HU-2011-1220	2011-01-01	2011-05-01	Second Class	AT-735	
3	11731	IT-2011-3647632	2011-01-01	2011-05-01	Second Class	EM-14140	
4	22255	IN-2011-47883	2011-01-01	2011-08-01	Standard Class	JH-15985	

	Customer Name	Segment	City	State	...	\
0	Toby Braunhardt	Consumer	Constantine	Constantine	...	
1	Joseph Holt	Consumer	Wagga Wagga	New South Wales	...	
2	Annie Thurman	Consumer	Budapest	Budapest	...	
3	Eugene Moren	Home Office	Stockholm	Stockholm	...	
4	Joseph Holt	Consumer	Wagga Wagga	New South Wales	...	

	Region_Central Asia	Region_EMEA	Region_East	Region_North	Region_North Asia	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	1	0	0	0	
3	0	0	0	1	0	
4	0	0	0	0	0	

	Region_Oceania	Region_South	Region_Southeast Asia	Region_West	profit
0	0	0	0	0	-105.437957
1	1	0	0	0	-35.829578
2	0	0	0	0	-29.526952
3	0	0	0	0	26.131462
4	1	0	0	0	-37.575104

[5 rows x 39 columns]

```
In [52]: from sklearn.model_selection import train_test_split

X = df.drop('Sales', axis=1) # Features
y = df['Sales'] # Target variable

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=
```

In [ ]: